

Softmax with temperature from the paper

"Distilling the Knowledge in a Neural Network"

Neural networks typically produce class probabilities by using a "softmax" output layer that converts the **logit, z_i** , computed for each class into a probability, q_i , by comparing z_i with the other logits.

$$q_i = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)} \quad (1)$$

2

where **T is a temperature** that is normally set to 1. Using a higher value for T produces a softer probability distribution over classes.

* Key Concept

Model learns **probability distribution**

Loss measures **difference between model distribution**
and data distribution

Optimizer trains model to **decrease loss**

Details about "Loss"

Training dataset $D = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$

Key assumption: independent and identically distributed (i.i.d.)

$$(x_i, y_i) \perp (x_j, y_j) \quad (x_i, y_i) \sim p(x, y)$$

$$\begin{aligned} \Rightarrow p(D) &= \prod p(x_i, y_i) \\ &= \prod p(x_i) p(y_i | x_i) \end{aligned}$$

$$p(D) = \prod p(x_i) p_\theta(y_i | x_i) \quad \text{model of true } p(y_i | x_i)$$

choose θ such that $p(D)$ is maximized

$$\log p(D) = \sum \cancel{\log p(x_i)} + \log p_\theta(y_i | x_i)$$

$$\theta^* \leftarrow \operatorname{argmax}_\theta \sum \log p_\theta(y_i | x_i) \quad \text{maximum likelihood estimation (MLE)}$$

$$\theta^* \leftarrow \operatorname{argmin}_\theta - \sum \log p_\theta(y_i | x_i) \quad \text{negative log-likelihood (NLL)} \\ \text{our loss function}$$

Information of an event E

$$I(E) = -\log_2 p(E)$$

Entropy of random variable X

$$H(X) = E(I(X))$$

$$= -\sum p(x_i) \log p(x_i)$$

Cross entropy of the distribution q relative to a distribution p

$$\begin{aligned} H(p, q) &= -E_p(\log q) \\ &= -\sum p(x) \log q(x) \end{aligned}$$

KL Divergence (Kullback - Leibler divergence)

relative entropy

statistical distance measuring how probability distribution Q is different from reference probability distribution P

model
data

$$\begin{aligned} D_{KL}(P \parallel Q) &= \sum p(x) \log \frac{p(x)}{q(x)} \\ &= -\sum p(x) \log \frac{q(x)}{p(x)} \end{aligned}$$

JS Divergence (Jensen - Shannon divergence)

$$JSD(P \parallel Q) = \frac{1}{2} D_{KL}(P \parallel \frac{P+Q}{2}) + \frac{1}{2} D_{KL}(Q \parallel \frac{P+Q}{2})$$

Summary

for 1 data point (x_i, y_i)

- cross entropy : $-\sum_y p(y|x_i) \log p_\theta(y|x_i)$
- negative log-likelihood : $-\log p_\theta(y_i|x_i)$
- KL divergence : $-\sum_y p(y|x_i) \log \frac{p_\theta(y|x_i)}{p(y|x_i)}$
- binary cross entropy : $-p(y|x_i) \log p_\theta(y|x_i) - (1-p(y|x_i)) \log (1-p_\theta(y|x_i))$

$$\text{KL Divergence} \longleftrightarrow \text{Cross Entropy}$$


$$-\sum p(x) \log q(x) + \sum p(x) \log p(x) \quad -\sum p(x) \log q(x)$$

$$\text{Cross Entropy} \longleftrightarrow \text{NLL}$$

Consider only $(x_i, y_i) \rightarrow$ model output is probability distribution

$$H(p, p_\theta) = -\sum_y p(y|x_i) \log p_\theta(y|x_i) \quad -\log p_\theta(y_i|x_i)$$

if $p(y|x_i)$ is one-hot encoded ...



$$\begin{aligned} H(p, p_\theta) &= -\sum_y p(y|x_i) \log p_\theta(y|x_i) \\ &= -p(y_1|x_i) \log p_\theta(y_1|x_i) - p(y_2|x_i) \log p_\theta(y_2|x_i) \\ &\quad \dots - p(y_n|x_i) \log p_\theta(y_n|x_i) \\ &= -p(y_i|x_i) \log p_\theta(y_i|x_i) \\ &= -\log p_\theta(y_i|x_i) \\ &= \text{NLL} \end{aligned}$$